



HT does not rescue MIP under skewed zero-mean labels: both become mode-seeking

TRAINED MODELS — 8 seeds/cell; HT = repository multivariate Student-t NLL ($\nu=2$, learned state scale); HG is its Gaussian control


A Skew makes failure reproducible for both HT and MIP




B HT and MIP hit the predicted -0.1 mode; HG keeps the mean



C The shared mode bias becomes a two-millimeter control miss



D Same recovery-covered data, different deployed fields



Skew cell (8 seeds): L2 8/8, HG 7/8, HT 0/8, MIP 0/8; bias HT -0.101 ± 0.001 , MIP -0.105 ± 0.004 (population prediction -0.100). Symmetry controls return both biases to ≈ 0 .