



Winner reversal: HT selects the clean branch while L2 and MIP follow contamination

TRAINED 2D CONTROL MODELS — 8 seeds; 30k recovery-covered chunks; 70% clean labels + 20% mild and 10% gross same-direction corruption

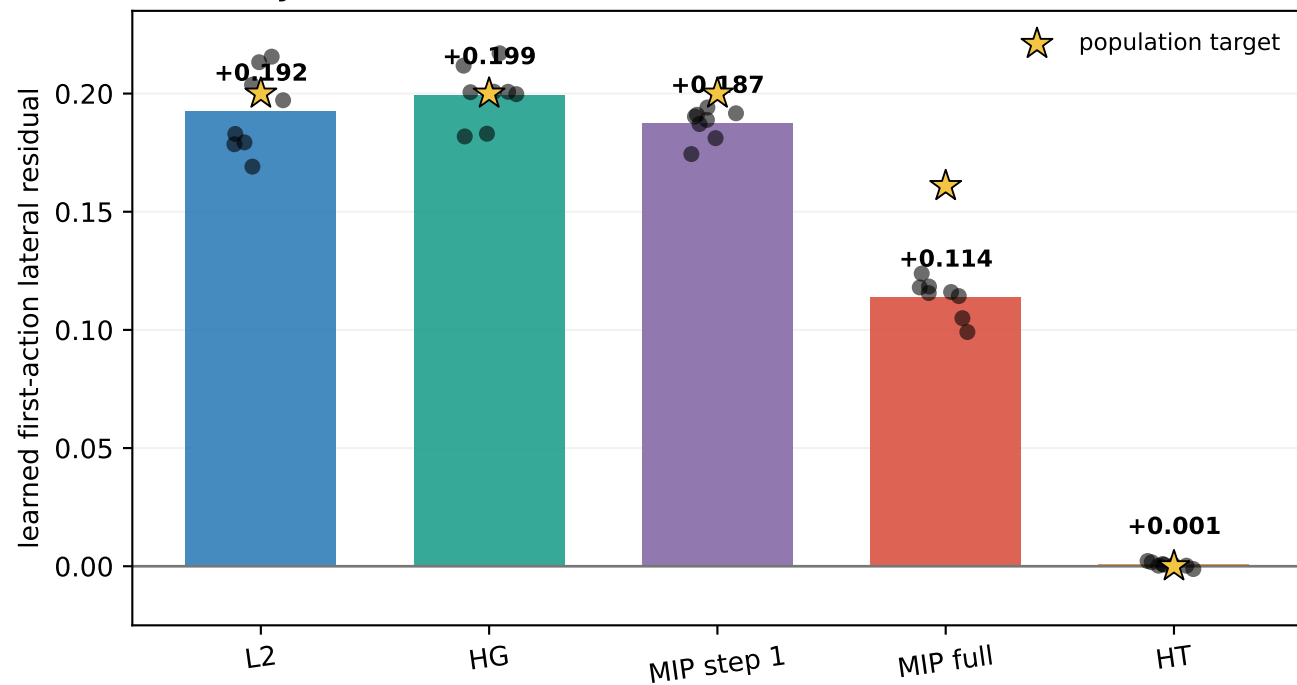
A Put the label mean on a corrupt branch



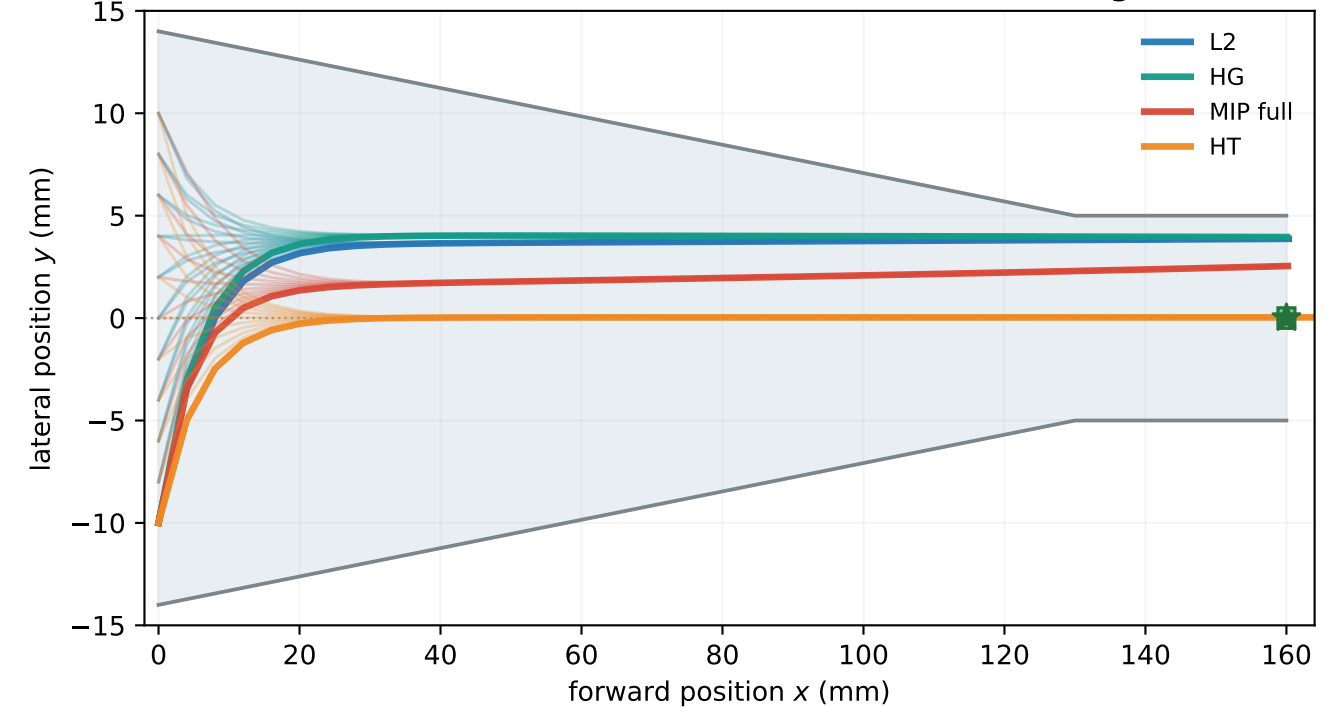
B HT is the only successful estimator



C HT rejects both contaminant scales and recovers zero



D Robust branch commitment is correct in this data regime



8 seeds: HT bias $+0.0007 \pm 0.0010$, dock error 0.02 mm; L2 3.91 mm; full MIP 2.50 mm. Same HT mode-seeking inductive bias as before —opposite outcome because the majority mode is task-correct.