Executed noisy paths remove the binary reversal: MIP and HT both recover precise control

TRAINED MODELS — 8 seeds; 3,000 dynamically valid paths / seed; 120k H=8 chunks; every noisy demonstration succeeds

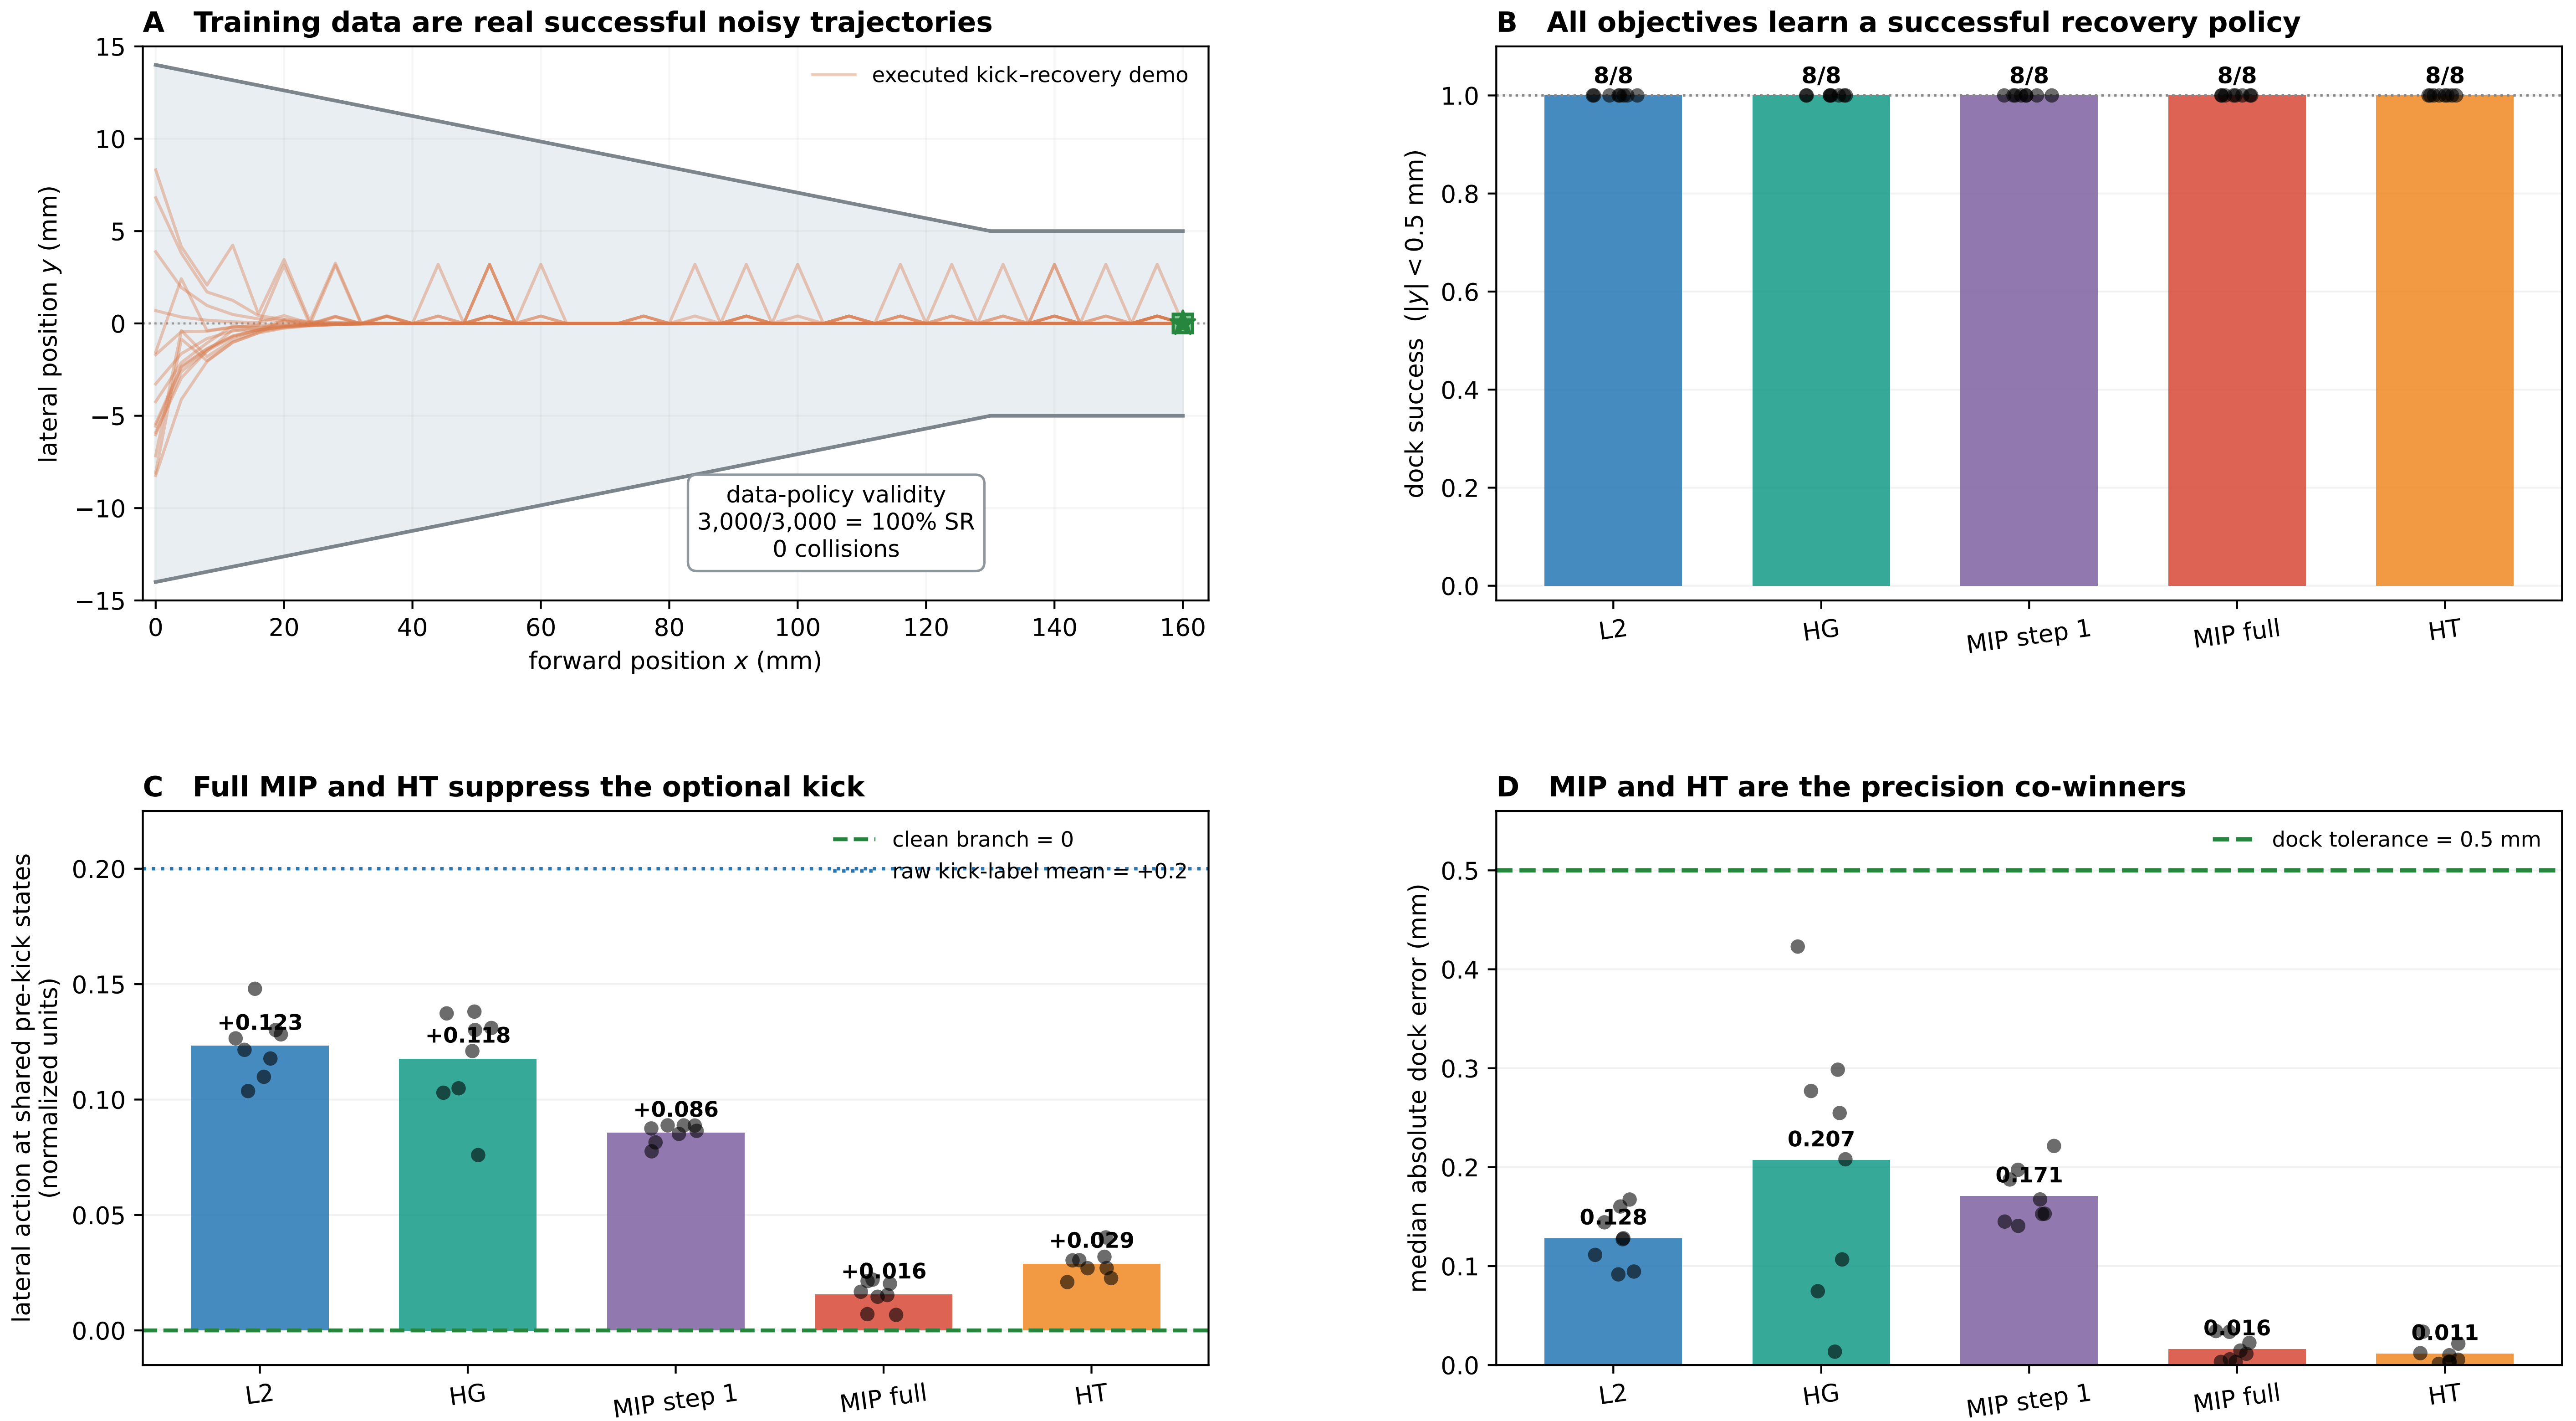

Eight-seed mean dock error: L2 0.128 mm, full MIP 0.016 mm, HT 0.011 mm. Full MIP is 8.0× more precise than L2, but SR is tied because every method remains inside the 0.5-mm dock.