


Diffusion on the obstacle toy: replanning without memory flips the branch decision; conditioning on the previous chunk fixes it

memoryless DP (resample each replan)

SR 0.248 ± 0.03 collision 0.70



DP + memory (previous chunk as input)

SR 0.932 ± 0.03 collision 0.05

