

One dataset, one training per method — the endpoint position alone decides the winner
(branch aims 0.8/0.2 at $\mp 6/+6$ mm + symmetric smear; 8 seeds, shading = seed std)

