2D no-free-lunch result: MIP's second-step denoiser turns skewed zero-mean nuisance into control bias

TRAINED MODELS — 8 seeds per cell, identical regression/MIP network class, 30k recovery-covered state-action chunks, H=8

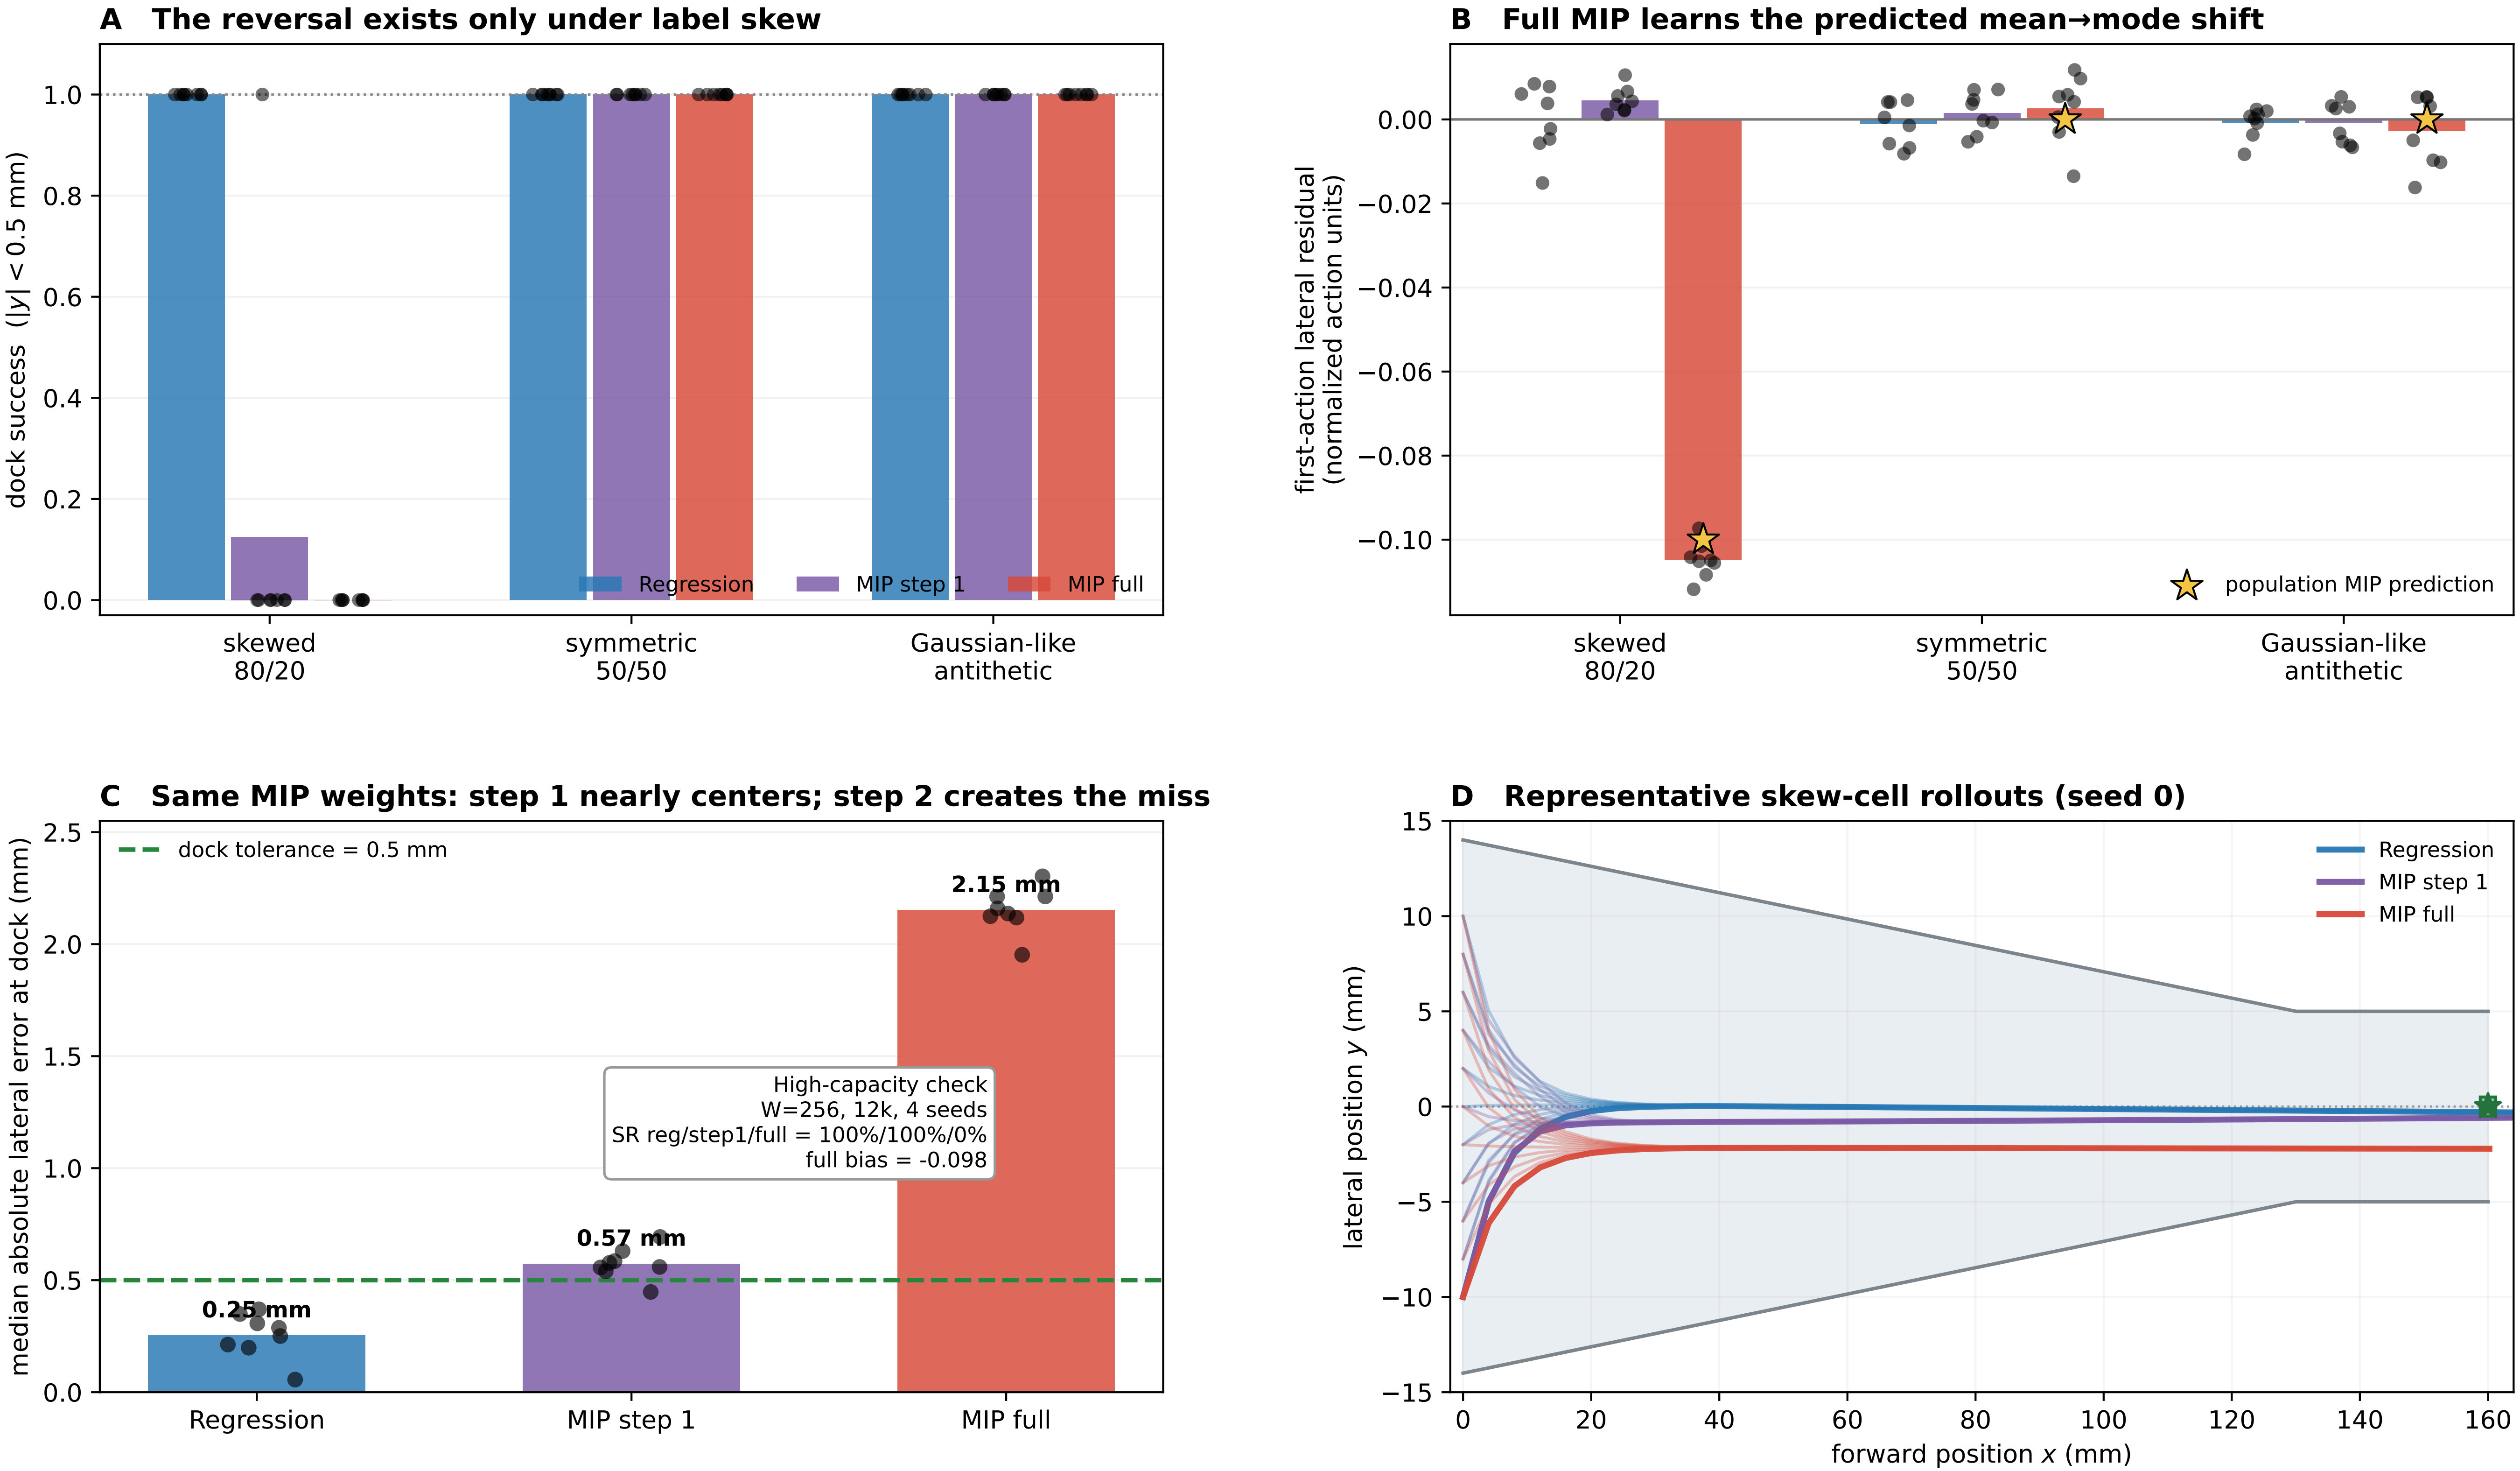

Skew cell: full-MIP bias -0.105±0.004 (analytic −0.100); dock error Regression 0.25, MIP-step1 0.57, full MIP 2.15 mm. Symmetric and Gaussian-like controls: 100% for all methods in all seeds.