


Collected data (top) vs trained policies (bottom) on the original skew task

A Collected demonstrations: corrective servo + executed chunk-held toyskew jitter (mostly at -2 with +8 excursions; zero-mean labels)



B Trained policies (seed 0): the mean family converges to 0; MIP full and HT converge to -2 — each endpoint rewards one family

