


Collected data (top) vs trained policies (bottom): wide-majority skew (b=0.2, sigma=0.1) separates all three

A Collected demonstrations: majority draws $-b + N(0, \sigma)$ (a BAND around -4 mm), minority $+4b$ kicks; zero-mean labels



B Trained policies (seed 0): L2 to the mean (0), MIP full to the posterior edge (~ -2), HT to the mode (~ -4) — three distinguishable bundles

