


Trained result on the original skew task: moving the endpoint from 0 to -2 mm reverses every verdict

TRAINED MODELS — 8 seeds per arm, original panel-A geometry (tight 0.5 mm tolerance), original toyskew nuisance {-b:.8, +4b:.2}; only the graded endpoint moves



Landings (median over seeds): Regression +0.06 / HG +0.10 / step1 +0.02 / MIP full -1.95 / HT nu=2 -1.99 / HT nu=0.5 -1.97 mm | the original witness graded only the left group of panel A; the right group is the same checkpoints at d = -2